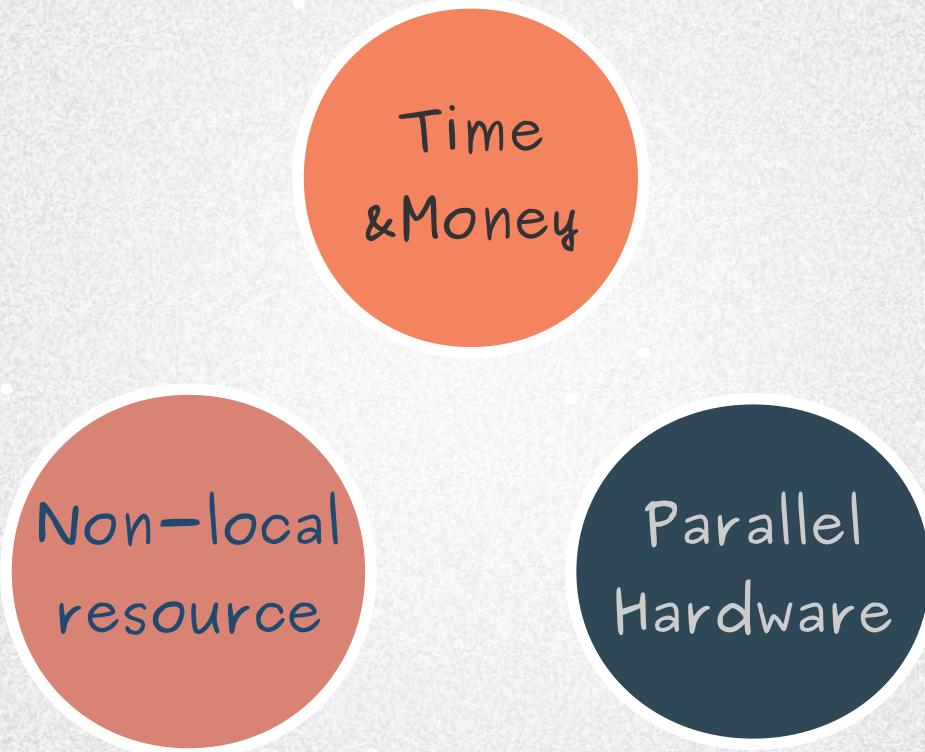


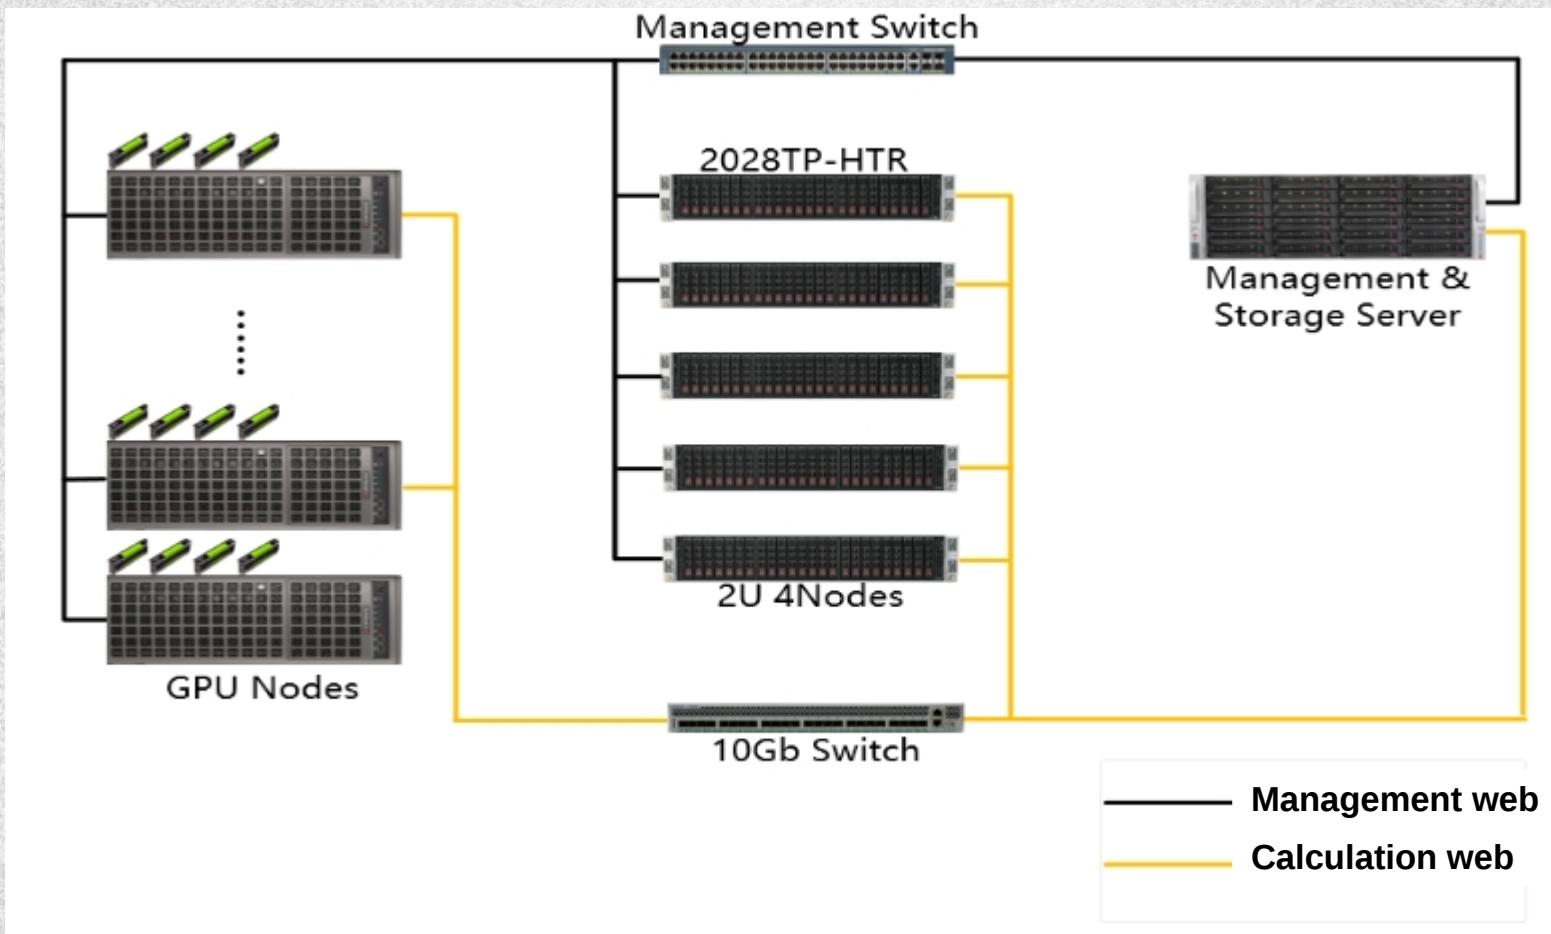
Parallel programming in Madagascar

Chenlong Wang

Why parallel ?



HPC structure



Outline

Parallel calculation in Madagascar with

- OpenMP
- MPI
- Pscons

OpenMp

Installation

OpenMP is a feature of the compiler and its parallel calculation is based on the **shared-memory**

In Madagascar, you can do parallel programming with the OpenMP either internally or externally

- **#include <omp.h> plus -fopenmp**
- **sfomp**

Internal usage

hello.c:

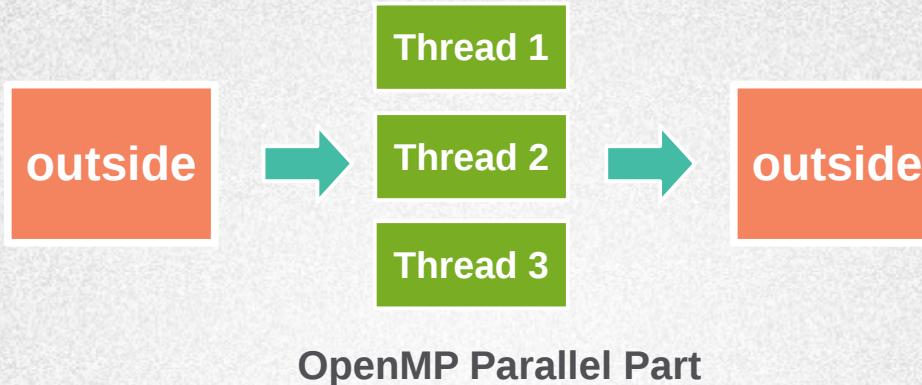
```
#include <omp.h>
#include <stdio.h>

int main()
{
    #pragma omp parallel
    {
        printf("Hello world from %d of %d\n", omp_get_thread_num(), omp_get_num_threads());
    }
    return 0;
}
```

compile

```
gcc hello.c -o hello -fopenmp / icc hello.c -o hello -openmp
```

Tutorials



The default scoping of variable is **shared** unless specified

- **#pragma omp parallel {}**
- **#pragma omp parallel for**
- **Reduction**

Exercise 1

```
#include <omp.h>
#include <stdio.h>

int main()
{
    #pragma omp parallel
    {
        #pragma omp master
        {
            printf("Hello world from %d \n", omp_get_thread_num());
        }

        #pragma omp single // Threads will meet barrier automatically
        {
            printf("Hello world from %d \n", omp_get_thread_num());
        }

        #pragma omp barrier
        printf("Number of threads %d \n", omp_get_num_threads());
    }
}
```

Code inside a parallel region will be copied to several copies and executed by different threads

Only responsible for master thread

Only be executed by a single thread

All threads will stop here until the slowest one finished

Exercise 2

gcc sin.c -fopenmp -lm -o sin

With OpenMP

Time 0.00328

Without OpenMP

Time 0.01896

```
#include <omp.h>
#include <stdio.h>
#include <math.h>

int main ()
{
    long i;
    long n=100000;

    double X[n],Y[n];

    for (i=0;i<n;i++)
        X[i]=(double)i;

    double wstart = omp_get_wtime();

    #pragma omp parallel for
    for (i=0;i<n;i++)
        Y[i]= sin(X[i])+cos(X[i]);  
    }  
    printf("Time %f\n",omp_get_wtime()-wstart);
```

i is private variable here

Exercise 3

$$\pi = \int_0^1 \frac{4.0}{1+x^2} dx$$

reduction

With Reduction

Time 0.073934

Without Reduction

Time 2.564278

```
#include <omp.h>
#include <stdio.h>

int main()
{
    long i;
    double x, mypi,sum=0.0;
    double vpi = 3.141592654;
    double tol = 0.0000001;
    long nsteps = 10000000;
    double step,diff,wend;

    double wstart = omp_get_wtime();

    step = 1.0/ (double) nsteps;

    // #pragma omp parallel for private (x) reduction(+:sum)
    #pragma omp parallel for private (x)
    for(i=0;i<nsteps;i++)
    {
        x = (i + 0.5) * step;
        #pragma omp atomic
        sum += 4.0/(1.0 + x * x);
    }

    mypi = step * sum;
    wend = omp_get_wtime();

    diff = mypi-vpi;
    if (diff<0) diff *= -1.0;

    if (diff>tol)
        printf("Error in pi: %f\n",mypi);
    else
        printf("PI %1.10f\n", mypi);
    printf("Time: %f\n",wend-wstart );
}
```

set x as private

execute summation successively

In Madagascar

```
grep "pragma omp" $RSFSRC/*/*/*/*.c |  
awk -F ':' '{ print $1 }' |  
uniq
```

139 standalone programs (approximately 11% of Madagascar programs) were using OMP on the last check (2014-02-09)

external usage

If the input data is supposed to be parallel, **sfomp** command splits the data along the given axis and runs it through parallel threads

OpenMP wrapper for embarrassingly parallel jobs.

`sfomp < inp.rsf > out.rsf split=ndim join=axis`

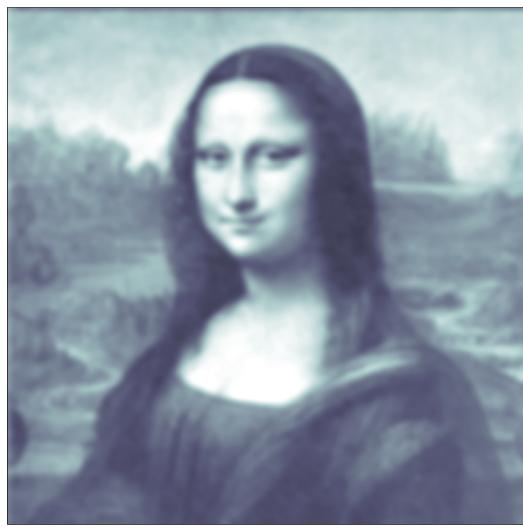
<i>int</i>	join=axis		axis to join
<i>int</i>	split=ndim		axis to split

`sfomp sfsmooth rect1=5 rect2=5 < in.rsf > out.rsf`

Smoothing Mona



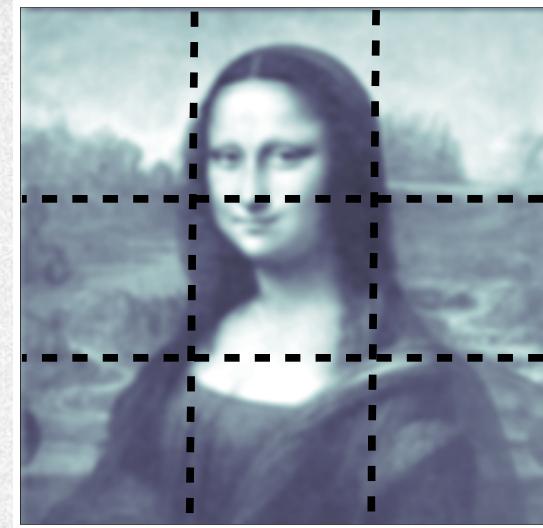
Mona Lisa



Smoothed without OpenMP

Mona Lisa

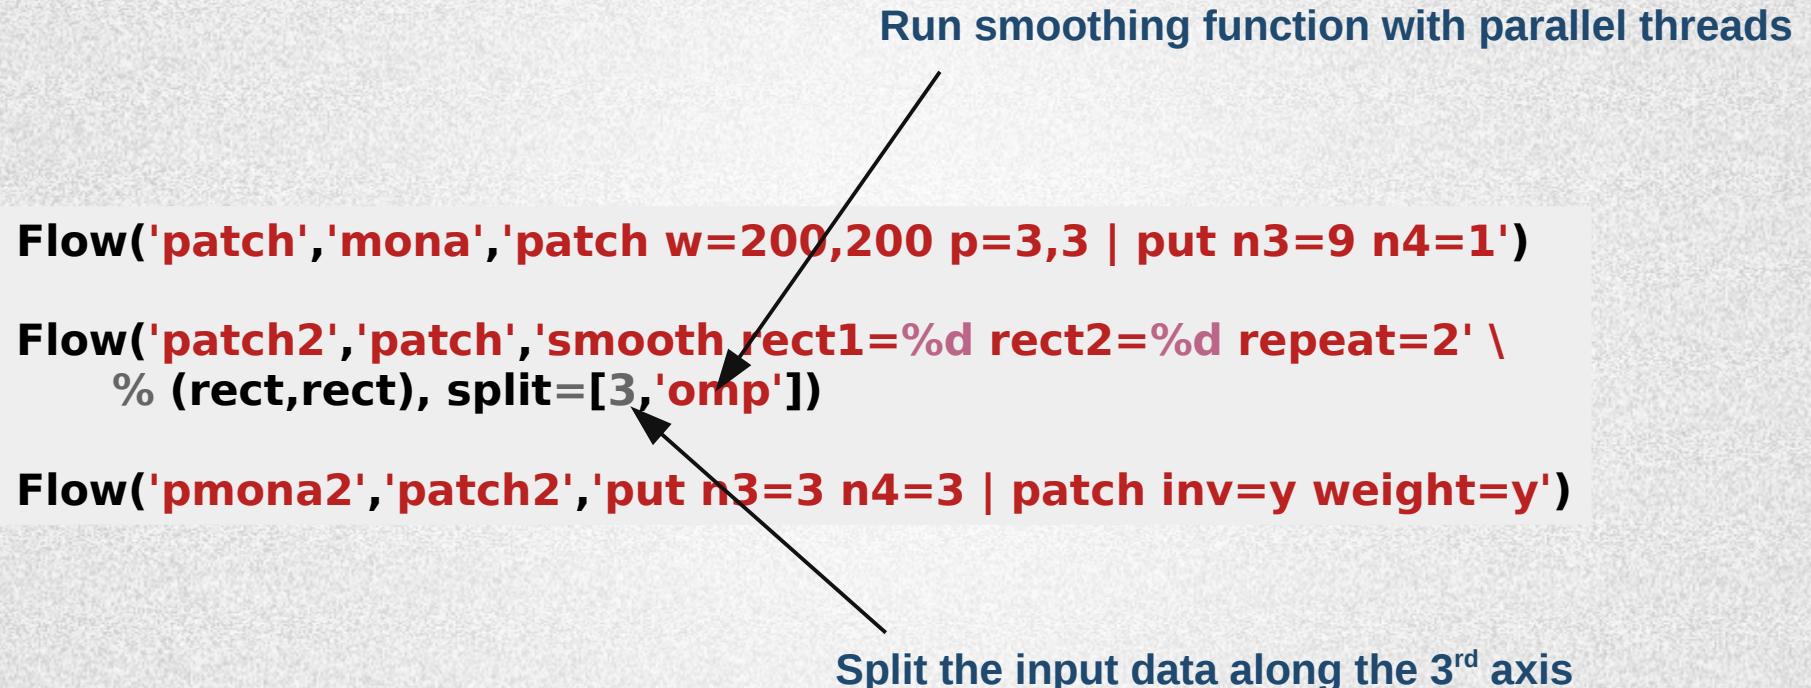
Without OpenMP



Smoothed with OpenMP

With OpenMP

Smoothing Mona



MPI

Installation

MPI (Message-Passing Interface) is dominant framework for parallel processing including **distributed-memory system. Several implementations (such as Open MPI and MPICH are available)**

In Madagascar, you can do parallel programming with the MPI either internally or externally as well

Internal usage

Mmpihelloworld.c:

```
#include <mpi.h>
#include <rsf.h>
#include <stdio.h>
#include <math.h>

int main(int argc, char *argv[])
{
    int myid, numprocs;
    int namelen;
    char processor_name[MPI_MAX_PROCESSOR_NAME];

    sf_init (argc, argv);

    MPI_Init(&argc,&argv);
    MPI_Comm_rank(MPI_COMM_WORLD,&myid);
    MPI_Comm_size(MPI_COMM_WORLD,&numprocs);
    MPI_Get_processor_name(processor_name,&namelen);

    fprintf(stderr,"Hello World! Process %d of %d on %s\n", myid, numprocs, processor_name);

    MPI_Finalize();
    exit(0);
}
```

The name starts with mpi

```
from rsf.proj import *

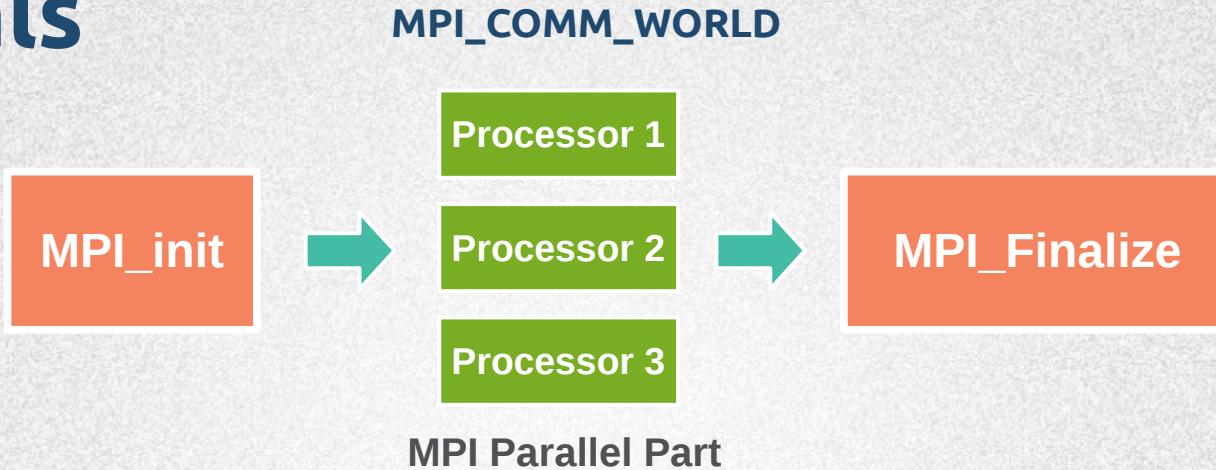
if not Wherels('mpirun'):
    sys.stderr.write("\nNo MPI.\n")
    sys.exit(1)

NP = int(ARGUMENTS.get('NP','4'))

Flow('out',None,
     '',
     "sfmpihelloworld -hostfile=hostfile",
     "'",np=NP,stdin=0,stdout=-1)

End()
```

Tutorials



MPI_COMM_WORLD is a handle points to all the resource

- **MPI_Send** and **MPI_Recv**
- **MPI_Datatype**

Exercise 1

Mmpihellocommu.c:

Send a message from master to slave

```
#include <mpi.h>
#include <rsf.h>
#include <stdio.h>
#include <string.h>
#include <math.h>

int main(int argc, char *argv[])
{
    int myid;
    char message[30];
    MPI_Status mpi_status;

    sf_init (argc, argv);
    MPI_Init(&argc,&argv);
    MPI_Comm_rank(MPI_COMM_WORLD,&myid);
    if (myid==0)
    {
        strcpy(message,"Hello processor 1!");
        MPI_Send(message, strlen(message), MPI_CHAR, 1, 99, \
                 MPI_COMM_WORLD);
    }
    else if(myid == 1)
    {
        MPI_Recv(message, 30, MPI_CHAR, 0, 99, \
                 MPI_COMM_WORLD, &mpi_status);
        fprintf(stderr, "received: %s\n", message);
    }

    MPI_Finalize();
    exit(0);
}
```

Exercise 2

Mmpihello.c:

```
/* Input vectors in memory */
.....
a = sf_floatalloc (n1);
b = sf_floatalloc (n1);
/* How many vectors per CPU */
nc = (int)(n2/(float)ncpu + 0.5f);
c = sf_floatalloc2 (n1, nc);
/* Starting position in input files */
sf_seek (ain, n1*cpuid*esize, SEEK_CUR);
sf_seek (bin, n1*cpuid*esize, SEEK_CUR);
for (i = cpuid; i < n2; i += ncpu, k++) {
    /* Read local portion of input data */
    sf_floatread (a, n1, ain);
    sf_floatread (b, n1, bin);
    /* Parallel summation here */
    for (j = 0; j < n1; j++)
        c[k][j] = a[j] + b[j];
    /* Move on to the next portion */
    sf_seek (ain, n1*(ncpu - 1)*esize, SEEK_CUR);
    sf_seek (bin, n1*(ncpu - 1)*esize, SEEK_CUR);
}
```

Splits the data along
the slowest axis

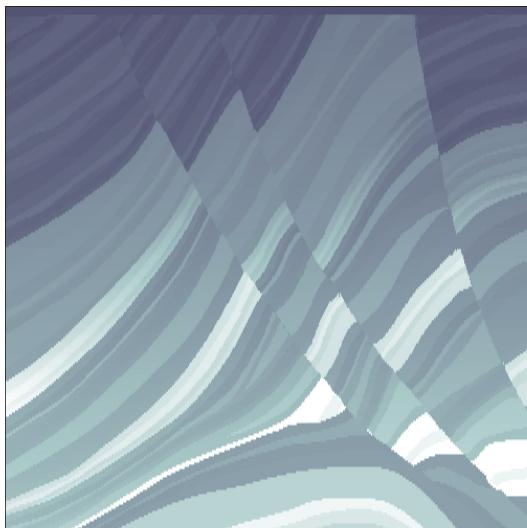
```
if (0 == cpuid) { /* Collect results from all nodes */
    for (i = 0; i < n2; i++) {
        k = i / ncpu; /* Iteration number */
        j = i % ncpu; /* CPU number to receive from */
        if (j) /* Receive from non-zero CPU */
            MPI_Recv (&c[k][0], n1, MPI_FLOAT, j, j,
                      MPI_COMM_WORLD, &mpi_stat);
        sf_floatwrite (c[k], n1, cout);
    }
    sf_fileclose (cout);
} else { /* Send results to CPU #0 */
    for (i = 0; i < k; i++) /* Vector by vector */
        MPI_Send (&c[i][0], n1, MPI_FLOAT, 0, cpuid,
                  MPI_COMM_WORLD);
}
sf_fileclose (ain); sf_fileclose (bin);
MPI_Finalize ();
return 0;
```

Collect results

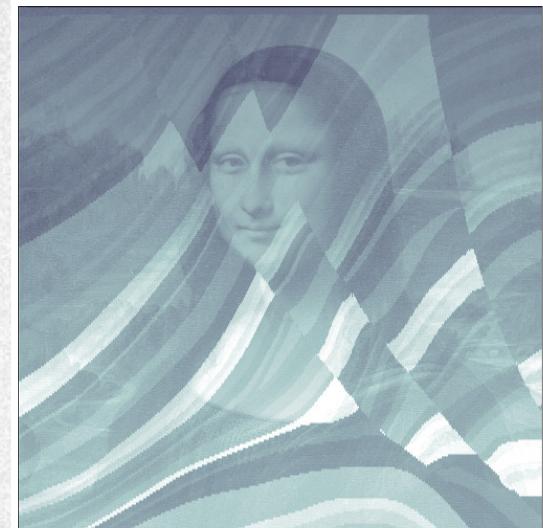
Exercise 2



Mona Lisa



Marmousi



Monamousi

Mona Lisa



Marmousi

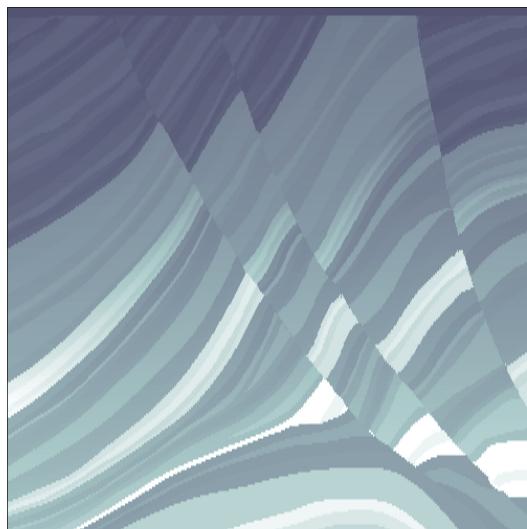


Monamousi

Exercise 2

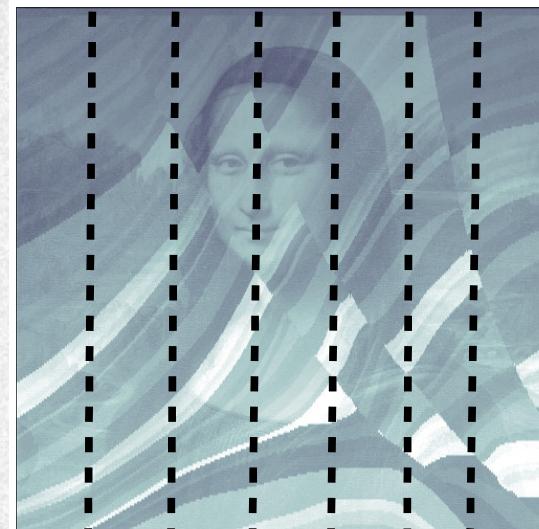


Mona Lisa



Marmousi

np



Monamousi

Mona Lisa



Marmousi



Monamousi

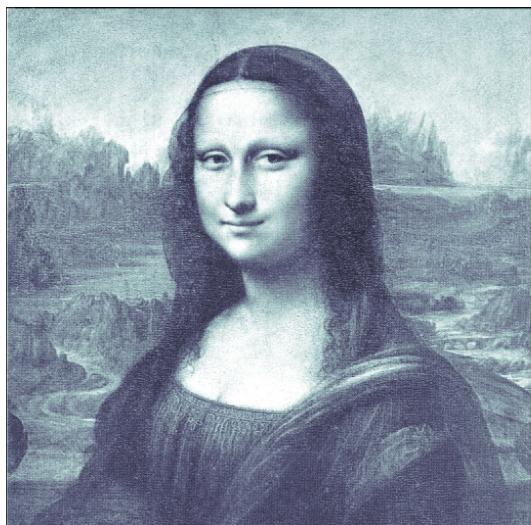
external usage

If the input data is supposed to be parallel, **smpi** command splits the data along the given axis and runs it through parallel threads

MPI wrapper for embarrassingly parallel jobs.		
smpi input=inp.rsf output=out.rsf split=ndim join=axis		
<i>string</i>	input=	auxiliary input file name
<i>int</i>	join=axis	axis to join
<i>file</i>	output=	auxiliary output file name
<i>int</i>	split=ndim	axis to split

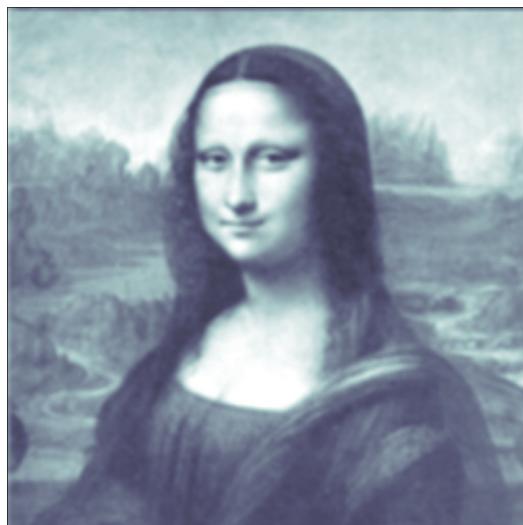
```
mpirun -np 4 smpi split=3 sfsmooth rect1=5 rect2=5 input=in.rsf output=out.rsf
```

Smoothing Mona



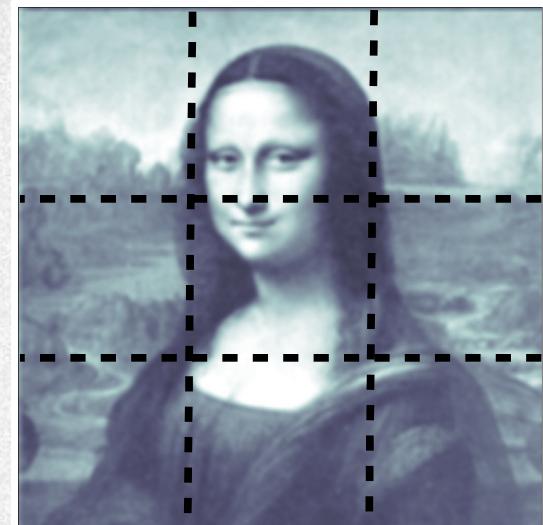
Mona Lisa

Mona Lisa



Smoothed without MPI

Without MPI



Patched and smoothed

With MPI

Smoothing Mona

Run smoothing function with parallel threads

```
Flow('patch','mona','patch w=200,200 p=3,3 | put n3=9 n4=1')
```

```
Flow('patch2','patch','smooth rect1=%d rect2=%d repeat=2' \
    % (rect,rect), split=[3,'mpi',[0]],reduce='cat',np=4)
```

```
Flow('pmona2','patch2','put n3=3 n4=3 | patch inv=y weight=y')
```

Split the input data along the 3rd axis

How to reduce (add or cat axis=1)

How many
processors
involved

Pscons

Pscons

The SCons can execute the script in parallel by splitting input data or splitting commands

Unlike the OpenMP or MPI utilities, this has fault tolerance -- in case of a node failing, restarting the job will allow it to complete.

Try

```
pscons / scons -j num
```

Demo1 – splitting input data

```
from rsf.proj import *

Fetch('mona.img','imgs')
Flow('mona','mona.img',
  ""
  echo n1=512 n2=513 in=$SOURCE data_format=native_uchar | dd type=float |
  window n1=512 n2=512 | sftransp | math output="input/1000"
  "",stdin=0)
Result('mona', ""grey allpos=y title="Mona Lisa" color=b screenratio=1 wantaxis=n"")

Fetch('marmvel.hh','marm')
Flow('mar','marmvel.hh',
  ""
  dd form=native | window j1=1 j2=2 f2=800 n2=512 f1=0 n1=512|
  scale dscale=0.001 |
  put label1=Depth unit1=km label2=Distance unit2=km | math output="input/4.45"
  "")
Result('mar', ""grey allpos=y title="Marmousi" color=b screenratio=1 wantaxis=n"")

Flow('monamousi','mar mona','math b=${SOURCES[1]} output="input+b"', split=[2,256])
Result('monamousi', ""grey allpos=y title="Monamousi" color=b screenratio=1 wantaxis=n"")

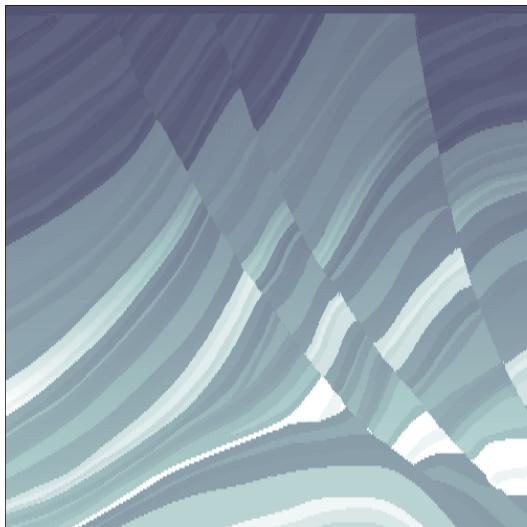
End()
```

The 2nd axis
Length 256

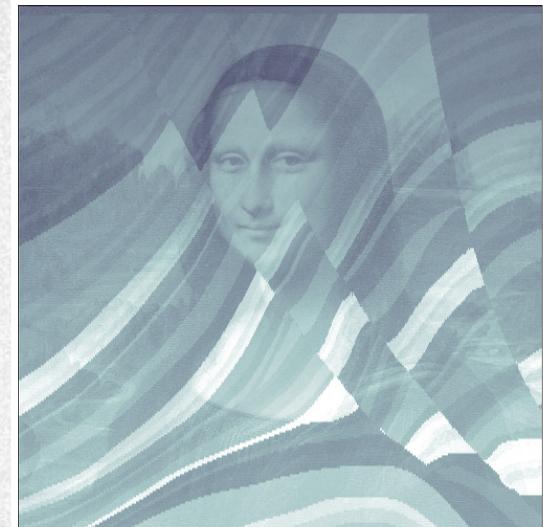
Demo 1



Mona Lisa



Marmousi



Monamousi

Mona Lisa



Marmousi

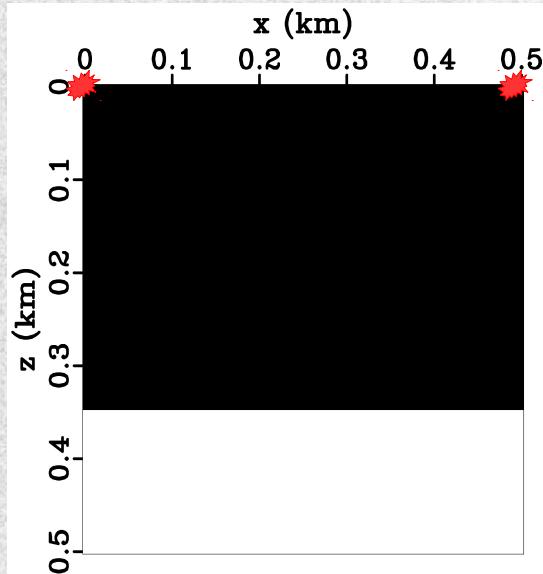


Monamousi

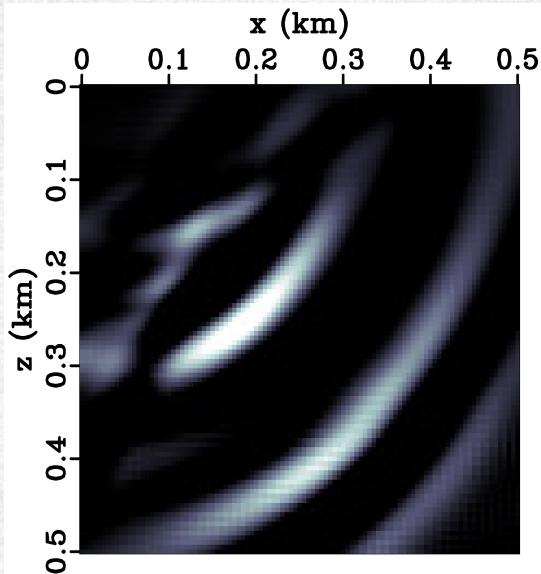
Demo2 – splitting commands

```
from rsf.proj import *
.....
#shot configuration in coordinate
nshot=2
shot0=0
dshot=0.5
Flow('sz',None,'math n1=1 output=%(sz)g' %par)
for i in range (1,nshot+1):
    isx = shot0 + (i-1)*dshot
    print isx
    Flow('sx-%d'%(i),None,'math n1=1 output=%g'%(isx))
    Flow('source-%d'%(i),'sx-%d sz'%(i),'cat axis=1 ${SOURCES[1]}')
    Flow('erecfield-%d ewavefield-%d'%(i,i),'elasticwavelet density receivers source-%d ccc'%(i),
         ...
         ewefd2d
         den=${SOURCES[1]}
         rec=${SOURCES[2]}
         sou=${SOURCES[3]}
         ccc=${SOURCES[4]}
         wfl=${TARGETS[1]}
         dabc=y snap=y verb=y jsnap=%(jsnap)d jdata=%(jdata)d
         ssou=y nb=20 nbell=11
         ... %par)
End()
```

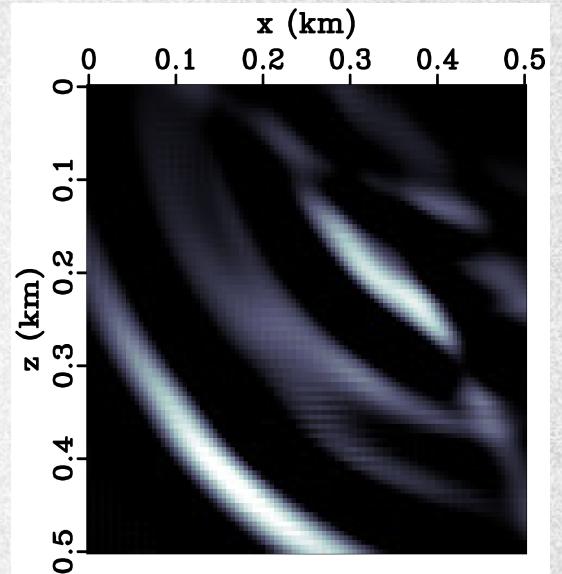
Demo 2



Two shots



Shot 1



Shot 2

Acknowledge

I would like to acknowledge the HPC group in NTNU gave me lectures on parallel programming ([https://www.hpc.ntnu.no/display/hpc/NTNU +HPC+GROUP](https://www.hpc.ntnu.no/display/hpc/NTNU+HPC+GROUP)).

I also appreciates the MSIM Research group provide some demos in MPI programming. (<https://www.msim.no>)